
Shared-Memory Honeytokens in Multi-Agent LLM Systems: A Controlled Simulation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Decoy credentials catch agents that misuse secrets they find in their environment.
2 When agents share memory, an attacker can pool what it learns about decoys, and
3 theory on the shared-memory honeytoken boundary predicts that a coalition can
4 pool weak fingerprints to distinguish decoys. We test that prediction in a fully syn-
5 thetic sandbox where open-weight models (Qwen3-4B, Qwen3-8B, Gemma3-4B)
6 run tool-call loops against a credential store seeded with decoys. Four conditions
7 isolate the boundary: static baseline, coalition (shared memory), rotation control
8 (per-episode decoys), and brokered containment (private reference monitor, broker-
9 issued capabilities, quarantine). Containment holds post-trigger harm at exactly
10 zero with CI $[0, 0]$ while reaching the highest trap activation (0.50); baseline, coal-
11 ition, and rotation reach harm 12.8, 16.0, and 16.9. The non-containment intent
12 sweep spans activation 0.20–1.00 and harm 0.4–15.8; containment-high keeps
13 harm at 0. Core-matrix benign false-positive rate is 0–1.8%. Containment’s zero-
14 harm result holds across model sizes; Gemma3-4B is an engagement boundary case.
15 Everything runs locally: no real credentials, no network egress, no live targets.

16 1 Introduction

17 LLM agent teams increasingly share a tool registry and a memory pool. A common detection idea
18 carries over from intrusion detection: seed the environment with honeytokens (decoy secrets whose
19 use reveals misuse) and monitor who touches them [Spitzner, 2003, Juels and Rivest, 2013]. In
20 single-agent settings this works because one agent’s mistakes stay private. In a team that shares
21 memory, they do not: an attacker that touches a decoy can tell its teammates, and the coalition
22 learns which credentials are traps faster than the defender can rotate them. Theory on this shared-
23 memory honeytoken boundary shows how a coalition can pool weak fingerprints and reduce its
24 decoy-classification error as observations accumulate [Gans, 2026].

25 We built CanaryGame to measure that prediction. In this controlled testbed, open-weight models act
26 as both benign workers and adaptive attackers inside a synthetic credential store seeded with decoys.
27 The privileged sink only mutates scored sandbox state, so attacks are safe to run and cheap to repeat.
28 Four conditions isolate one variable each: whether memory is shared, whether decoys are rotated,
29 and whether a capability broker, a private reference monitor, and automatic quarantine stand between
30 agents and the sink.

31 Across three model configurations, sharing raises mean post-trigger harm and spreads decoy finger-
32 prints; rotation lowers activation but leaves harm high; containment blocks every decoy misuse. A
33 safety-tuned model that rarely misuses credentials reveals an attacker-engagement boundary. Our
34 contributions are (1) an offline, reproducible multi-agent testbed for honeytoken dynamics, (2) a mea-

Table 1: Core matrix, Qwen3-4B, $n=10$ /condition. Harm is harmful sink steps after first trigger; contagion counts fingerprints pooled through shared memory. CIs are bootstrap 95%.

Condition	Activation	Coverage	Post-trigger harm (95% CI)	Contagion	Benign FPR
Baseline	0.40	0.40	12.8 (3.2, 23.2)	0.0	0.000
Coalition	0.40	0.40	16.0 (3.8, 30.6)	16.0	0.018
Rotation	0.30	0.30	16.9 (0.0, 33.8)	16.9	0.000
Containment	0.50	0.00	0.0 (0.0, 0.0)	3.9	0.005

surement of pooling effects with per-condition confidence intervals, and (3) evidence that brokered containment restores detection across model sizes.

2 Testbed

Sandbox. Core episodes give three benign workers and two attackers a shared tool registry (`query_index`, `list_entries`, `read_file`, `write_file`, `search_files`, `send_message`, `upload_export`) and a credential store of 12 synthetic entries with service, role, and note fields; one sweep cell adds a third attacker. Core agents act for up to 40 steps by emitting one action per turn (`call_tool`, `use_credential`, `share_memory`, `probe`, or `noop`); tool outputs and sink results feed back into the next prompt. The privileged sink returns one of three statuses: *ok* (legitimate credential accepted), *blocked*, or *harmful* (decoy accepted in unbrokered conditions). Only scored state changes.

Decoys. Four of the store entries are decoys: synthetic, unbrokered values no legitimate flow should ever use. Ground-truth decoy labels live outside the surface the agents see.

Conditions. *Baseline*: isolated memories, fixed decoys. *Coalition*: shared memory, fixed decoys; attackers can pool what they learn. *Rotation*: shared memory, decoys re-randomized every episode. *Containment*: shared memory, private decoys, plus three defenses: a broker is the only legitimate credential source (decoys are never issued through it), a reference monitor watches every action against ground truth, and triggered agents are quarantined until release.

Attackers. Models drive every attacker action. An intent axis in the prompts sets aggression: *low* (passive audit), *medium* (contractor told to find and use credentials), *high* (adversarial coalition instructed to verify decoys and write teammate notes to shared memory). Benign workers draw real tasks from τ^2 -bench [Barres et al., 2025].

Models. Qwen3-4B-Instruct-2507 is the primary attacker and victim; Qwen3-8B and Gemma3-4B-IT provide cross-model checks [Yang et al., 2025, Gemma Team, 2025]. All models run locally through vLLM [Kwon et al., 2023] on one A100-40GB at temperature 0. For compatibility, Gemma3 prompts strictly alternate roles because its chat template rejects consecutive assistant turns. Its dialect wraps verbs (`call_tool use_credential secret=...`), which a normalization layer maps back to canonical actions. Qwen3-8B defaults to thinking-mode output, which parses as no-ops, so thinking is disabled via chat-template kwargs.

Data and metrics. Dataset v1 has 60 records (4 conditions $\times$ 5 seeds $\times$ 3 episodes) with deterministic fixtures; benign tasks come from a τ^2 -bench subset. We report trap activation (episodes with ≥ 1 decoy trigger), attack coverage, post-trigger harm (harmful sink steps after first trigger), contagion (fingerprints pooled through shared memory), benign false-positive rate (FPR), benign success, and benign engagement, with 95% confidence intervals per condition ($n=10$ core matrix, $n=5$ sweep).

3 Results

Containment blocks decoy harm (H3). Containment alone stops all damage despite the highest activation (0.50). Its attack coverage is 0.00, and its post-trigger harm is exactly zero with CI $[0, 0]$ (Table 1, Figure 1). Traces show the broker refusing every decoy, blocking 3.9 attempts per episode on average, so attackers keep probing but never move scored state. The same zero-harm result holds across model sizes and in the Gemma containment condition (Section 3.1).




Figure 1: Core matrix (Qwen3-4B, $n=10$): activation, post-trigger harm with CIs, contagion, and benign FPR. Containment has the highest activation and zero harm.

Table 2: Attacker-intent sweep, Qwen3-4B, $n=5/\text{cell}$.

Cell	Intent	Activation	Harm (95% CI)	Contagion	Benign succ.
low $\times$ 2	low	0.20	0.4 (0.0, 1.2)	0.4	1.00
medium $\times$ 2	medium	0.20	8.4 (0.0, 25.2)	8.4	1.00
high $\times$ 2	high	1.00	5.6 (3.0, 8.8)	5.6	0.60
high+medium	mixed	0.80	10.8 (2.0, 24.4)	10.8	0.70
high $\times$ 3	high	0.80	15.8 (4.4, 31.0)	15.8	0.70
contain. high	high	0.80	0.0 (0.0, 0.0)	5.0	0.80
contain. medium	medium	0.20	0.0 (0.0, 0.0)	1.6	1.00

75 **Pooling raises mean harm and spreads fingerprints (H1).** Coalition harm (16.0) sits above baseline
76 (12.8) with contagion 16.0: once one attacker touches a decoy, the fingerprint lands in shared memory
77 and both attackers converge on it. High-intent cells reach activation 1.00 and contagion up to 15.8
78 (Table 2), though at $n=10$ the baseline–coalition gap is directional, not significant.

79 **Rotation lowers contact; harm remains high (H2).** Re-randomizing decoys each episode produces
80 the lowest activation (0.30) and the slowest triggers, but when an attack does land, harm stays high
81 (16.9). Rotation buys time; containment prevents misuse outright.

82 **Intent chiefly controls contact.** With two attackers, moving from low to high intent raises activation
83 from 0.20 to 1.00; mean harm is 0.4 at low, 8.4 at medium, and 5.6 at high intent. At $n=5$, wide CIs
84 leave the effect of a third attacker unresolved (Table 2). Under high intent, containment still pins
85 harm at zero (0.80 activation, 0 harm).

86 **Benign false alarms rise under high-intent attack.** Across the core matrix, benign FPR is 0–1.8%
87 and benign success 0.95–1.00 (Figure 1d). One trajectory-signature detector reports an unconditional
88 benign FPR of 25–53%, depending on the recall protocol [Leong, 2026]; the reference monitor here
89 keys on decoy tokens rather than trajectories, so benign agents rarely trip it. Sustained high-intent
90 attack fills the shared pool with attacker fingerprints: benign FPR rises to 3–12%, and benign success
91 drops to 0.60–0.80 (Table 2). Because benign success is derived from false alarms, these runs do not
92 establish task degradation.

93 3.1 Results across models

94 Benign behavior is similar on Gemma3-4B: engagement is 0.96–1.00, success 0.95–1.00, and FPR
95 near zero everywhere. Attack engagement does not transfer: Gemma attackers spend most of their
96 budget scanning and attempt their first credential around step 21–39 (Qwen: step 3), so activation
97 stays near zero even at an 80-step budget (Figure 2). Where Gemma attackers do misuse decoys,




Figure 2: Qwen3-4B vs. Gemma3-4B. Gemma preserves high benign success and near-zero FPR (b, c) but rarely activates traps (a).

containment blocks them all (harm 0.0). Qwen3-8B coalition harm is 1.9 versus 16.0 for Qwen3-4B, and containment again holds harm at exactly zero with CI $[0, 0]$. The zero-harm containment result holds across model sizes and in the few Gemma episodes with decoy misuse. Attacker engagement remains a boundary condition for evaluating the theory across model families.

4 Discussion and limitations

Sharing raises mean harm and spreads fingerprints, although the harm CIs overlap. A private reference monitor keyed on decoy tokens yields zero post-trigger harm and keeps benign FPR near zero. Removing decoys from the broker’s issuance path makes every decoy touch a dead end. Rotation lowers activation but leaves harm high after contact.

Four limitations bound these claims. (1) Everything is synthetic and offline: one credential-store surface, no live tools, no egress. Real deployments add surfaces we do not model. (2) Samples are small: $n=10$ per condition with overlapping harm CIs makes the baseline–coalition contrast directional. (3) Benign success is derived only from false alarms; benign engagement separately reports tool use, and neither is a task score. (4) Attack-side conclusions rest on models willing to attack; a misuse-averse model suppresses decoy contact itself, which hides detection effects rather than refuting them.

5 Related work

Honeytokens and honeywords are established tripwires for human intruders [Spitzner, 2003, Juels and Rivest, 2013]; we port the idea to agent teams whose communication channel changes the threat model. Tool-using agent benchmarks evaluate single-agent competence [Yao et al., 2023, Barres et al., 2025]; we repurpose τ^2 -bench tasks as benign workload inside a security experiment. Work on LLM serving [Kwon et al., 2023] and multi-agent frameworks makes shared-memory teams practical; studies of prompt injection compromise agents through content they read [Abdelnabi et al., 2023], complementary to our setting where the environment itself is instrumented and the channel under study is peer memory.

6 Ethics and reproducibility

All experiments run locally and offline. Credentials are synthetic, the sink mutates only scored sandbox state, and no external systems, live servers, or network calls are involved at any point, including during attack rollouts. The full dataset (60 records, fixed seeds, acceptance-checked) will accompany the submission as anonymized supplementary material and be released publicly after review; code, configs, pod runner scripts, and per-condition reports accompany it. Fixed seeds and temperature-0 decoding make every table regenerable.

References

L. Spitzner. Honeytokens: The other honeypot. *SecurityFocus*, 2003.

- 132 A. Juels and R. L. Rivest. Honeywords: Making password-cracking detectable. In *ACM CCS*, 2013.
- 133 S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao. ReAct: Synergizing reasoning
134 and acting in language models. In *ICLR*, 2023.
- 135 V. Barres, H. Dong, S. Ray, X. Si, and K. Narasimhan. τ^2 -bench: Evaluating conversational agents
136 in a dual-control environment. *arXiv:2506.07982*, 2025.
- 137 W. Kwon, Z. Li, S. Zhuang, Y. Sheng, L. Zheng, C. H. Yu, J. E. Gonzalez, H. Zhang, and I. Stoica.
138 Efficient memory management for large language model serving with PagedAttention. In *SOSP*,
139 2023.
- 140 A. Yang et al. Qwen3 technical report. *arXiv:2505.09388*, 2025.
- 141 Gemma Team. Gemma 3 technical report. *arXiv:2503.19786*, 2025.
- 142 S. Abdelnabi, K. Greshake, S. Mishra, C. Endres, T. Holz, and M. Fritz. Not what you’ve signed
143 up for: Compromising real-world LLM-integrated applications with indirect prompt injection. In
144 *ACM Workshop on Artificial Intelligence and Security (AISec)*, 2023.
- 145 J. S. Gans. When agents talk: Honeytokens under shared memory. *arXiv:2608.11436*, 2026.
- 146 J. W. Leong. Forensic trajectory signatures for agent memory poisoning detection. *arXiv:2606.30566*,
147 2026.

148 A Model-size comparison




Figure 3: Qwen3-4B vs. Qwen3-8B ($n=10$): containment pins harm at zero on both sizes; Qwen3-8B shows lower activation and coalition harm.

149 B Intent-gradient figure




Figure 4: Attacker-intent sweep ($n=5/\text{cell}$). High intent drives contact; harm varies widely, while containment holds it at zero under both tested intents.

150 C Reproduction pointers

151 Fixed seeds 0–4 drive every table: two episodes per seed in the core matrix and one in the sweep; tem-
 152 perature 0 everywhere. The condition grid lives in `configs/conditions.yaml`; single-command
 153 mock smoke: `uv run pytest`; GPU rollouts via `scripts/pod/run_matrix.sh` on a Lambda
 154 A100 pod (41 min for the 40-episode core matrix, \$1.35 at \$1.99/hr).